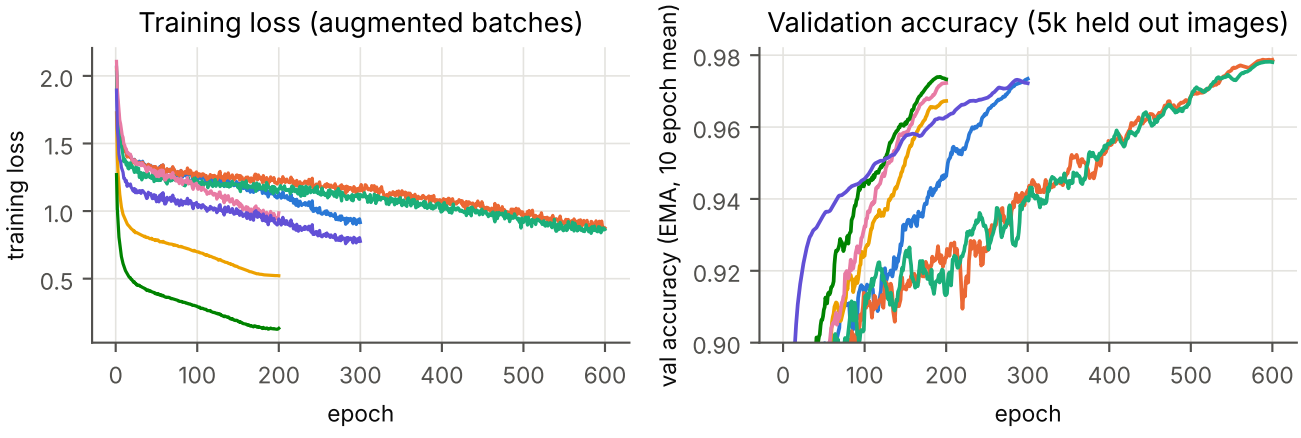
Training loss (augmented batches)
training loss
2.0
1.5
1.0
0.5
0
100
200
300
400
500
600
epoch
Validation accuracy (5k held out images)
val accuracy (EMA, 10 epoch mean)
0.98
0.96
0.94
0.92
0.90
0
100
200
300
400
500
600
epoch
v2 (ResNet-24, 300 ep)
v4 (SE-ResNet-26, 600 ep)
robust (v4, robust augmentation)
v1 (baseline, 200 ep)
SAM (200 ep)
distillation (200 ep)
AugMix (300 ep)